

Bad-Scoping Case (Short Version — Retail Bank)

A mid-sized retail bank is exploring a conversational AI assistant to support its loan approval teams as part of a broader digital transformation. The assistant would surface relevant applicant information, highlight key risk signals, and provide a recommended credit decision with a natural-language explanation inside the existing loan processing system. The initiative aims to improve turnaround time, increase decision consistency, and help onboard new credit analysts faster. The current proposal is to adapt a large language model using historical loan records, internal policy documents, and customer interaction data, with near real-time responses that fit into today's workflow. The program has executive sponsorship, access to enterprise data, and a target rollout within two quarters, with the expectation that it will meet regulatory and model governance standards.

Analysis

Sentence 1

“A mid-sized retail bank is exploring a conversational AI assistant to support its loan approval teams as part of a broader digital transformation.”

What it sounds like:

Clear business initiative with a defined user group.

Hidden scoping gap → Decision not defined

- Is the system making the approval decision?
- Recommending?
- Summarizing?
- Pre-screening?

These are completely different systems.

Strong learner should ask:

👉 *“What exact decision will this system support?”*

👉 *“Who is accountable for that decision?”*

Sentence 2

“The assistant would surface relevant applicant information, highlight key risk signals, and provide a recommended credit decision with a natural-language explanation inside the existing loan processing system.”

What it sounds like:

Well-scoped functionality.

Hidden scoping gaps → System function ambiguity

You now have **three different AI problems**:

1. Retrieval
2. Risk prediction
3. Decision explanation

Each requires:

- Different data
- Different evaluation
- Different governance

Strong learner should ask:

👉 *“Is this prediction, retrieval, or decision automation?”*

👉 *“Is the explanation regulatory or just internal?”*

Sentence 3

“The initiative aims to improve turnaround time, increase decision consistency, and help onboard new credit analysts faster.”

What it sounds like:

Clear success metrics.

Hidden scoping gap → No ML success criteria / no error cost

These are **business KPIs**, not system evaluation metrics.

Missing:

- Accuracy target?
- Recall for bad loans?
- Cost of false approval vs false rejection?

Strong learner should ask:

👉 *“What does ‘good’ look like?”*

👉 *“What is the cost of a wrong recommendation?”*

Sentence 4

“The current proposal is to adapt a large language model using historical loan records, internal policy documents, and customer interaction data, with near real-time responses that fit into today’s workflow.”

What it sounds like:

Data exists + architecture direction.

Hidden scoping gaps → Data reality + WHERE constraints

Data gaps:

- Where do labels for loan outcome come from?
- Are policies time-versioned?
- Is historical data biased?
- Is decision-time data the same as training data?

Operational gaps:

- What does “near real time” mean? (2 sec? 30 sec? 5 min?)
- Is the system in the approval loop or before it?

Strong learner should ask:

👉 *“Is this data available at decision time?”*

👉 *“Are historical approvals the ground truth — and are they biased?”*

Sentence 5

“The program has executive sponsorship, access to enterprise data, and a target rollout within two quarters, with the expectation that it will meet regulatory and model governance standards.”

What it sounds like:

Feasible and well-supported.

Hidden scoping gaps → Governance & risk undefined

Regulation in lending implies:

- Explainability requirements
- Adverse action reason tracking
- Fair lending bias monitoring
- Audit trail
- Human override accountability

None are defined.

Timeline pressure introduces:

- Scope risk
- Data readiness risk

Strong learner should ask:

👉 *“What level of explainability is required for regulatory review?”*

👉 *“Can a black-box recommendation even be used in this decision?”*

👉 *“Is human-in-the-loop mandatory?”*

BREAKTHROUGH LINE

“So this sounded like a well-scoped AI initiative.

But in reality:

The decision is unclear.

The data is unvalidated.

The risk is undefined.

And until those three are clear — there is no architecture to design.”
